

A graphic illustration featuring a dark background with a series of flowing, wavy lines in red and grey. On the left, a vertical column of binary code (0s and 1s) is visible. The right side of the image is filled with numerous lightbulb icons, some glowing red and others grey, connected by thin lines, suggesting a network or flow of ideas. The overall theme is technology and innovation.


22 de septiembre de 2026

22 de septiembre de 2026





Solución a un problema del milenio




 **OpenAI** 🌟 @OpenAI · 8 sept.

We're sharing a solution to the Navier-Stokes Millennium Prize Problem, one of the deepest problems at the frontier of mathematics.

The proof was produced by a group of agents, using an OpenAI next-generation model significantly more capable than GPT-6 Astra.

The problem concerns whether the description of smooth three-dimensional fluid motion modeled by the Navier-Stokes equations can break down. It has remained unresolved for roughly 90 years.



Luis Martínez-Zorúa (CUNEF Universidad) y Diego Córdoba (ICMAT-CSIC)



Tristan Buckmaster (NYU)
 "I believe Luis Martínez-Zorúa deserves a Fields Medal"



One
word
leads


Entrenamiento
no supervisado


the

next

bright

idea





Language Models are Few-Shot Learners

Tom B. Brown*	Benjamin Mann*	Nick Ryder*	Melanie Subbiah*
Jared Kaplan†	Prafulla Dhariwal	Arvind Neelakantan	Pranav Shyam
Girish Sastry	Amanda Askell	Sandhini Agarwal	Ariel Herbert-Voss
Gretchen Krueger	Tom Henighan	Rewon Child	Aditya Ramesh
Daniel M. Ziegler	Jeffrey Wu	Clemens Winter	Christopher Hesse
Mark Chen	Eric Sigler	Mateusz Litwin	Scott Gray
Benjamin Chess	Jack Clark	Christopher Berner	
Sam McCandlish	Alec Radford	Ilya Sutskever	Dario Amodei

OpenAI

Abstract

Recent work has demonstrated substantial gains on many NLP tasks and benchmarks by pre-training on a large corpus of text followed by fine-tuning on a specific task. While typically task-agnostic in architecture, this method still requires task-specific fine-tuning datasets of thousands or tens of thousands of examples. By contrast, humans can generally perform a new language task from only a few examples or from simple instructions – something which current NLP systems still largely struggle to do. Here we show that scaling up language models greatly improves task-agnostic, few-shot performance, sometimes even reaching competitiveness with prior state-of-the-art fine-tuning approaches. Specifically, we train GPT-3, an autoregressive language model with 175 billion parameters, 10x more than any previous non-sparse language model, and test its performance in the few-shot setting. For all tasks, GPT-3 is applied without any gradient updates or fine-tuning, with tasks and few-shot demonstrations specified purely via text interaction with the model. GPT-3 achieves strong performance on many NLP datasets, including translation, question-answering, and cloze tasks, as well as several tasks that require on-the-fly reasoning or domain adaptation, such as unscrambling words, using a novel word in a sentence, or performing 3-digit arithmetic. At the same time, we also identify some datasets where GPT-3’s few-shot learning still struggles, as well as some datasets where GPT-3 faces methodological issues related to training on large web corpora. Finally, we find that GPT-3 can generate samples of news articles which human evaluators have difficulty distinguishing from articles written by humans. We discuss broader societal impacts of this finding and of GPT-3 in general.

The three settings we explore for in-context learning

Zero-shot

The model predicts the answer given only a natural language description of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 cheese => ..... ← prompt
```

One-shot

In addition to the task description, the model sees a single example of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← example
3 cheese => ..... ← prompt
```

Few-shot

In addition to the task description, the model sees a few examples of the task. No gradient updates are performed.

```
1 Translate English to French: ← task description
2 sea otter => loutre de mer ← examples
3 peppermint => menthe poivrée ←
4 plush girafe => girafe peluche ←
5 cheese => ..... ← prompt
```

Reglas del mundo:


$$x_i^{t+1} = \begin{cases} 1 & \text{if } n_i^t = 3 \\ x_i^t & \text{if } n_i^t = 2 \\ 0 & \text{otherwise} \end{cases}$$

$$n_i^t = \sum_{j \in \mathcal{N}(i)} x_j^t$$




Simuladores


Juego de la vida
de Conway




Still life




Spaceship
(glider)





Oscillator
(pulsar)



Turing Machine




Simuladores
Juego de la vida
de Conway




Reglas simples



Universal
+
Irreducible



Simuladores
LLM




Regla básica del LLM: predecir siguiente token en base a los anteriores. Es aprendida estadísticamente.

En el límite, si un LLM contara con datos ilimitados de entrenamiento y convergiera a error 0 de predicción de tokens, se convertiría en un **simulador perfecto del "mundo en texto"**.

Ortogonalidad: el LLM no tiene agencia, no tiene objetivos, ni toma decisiones.

Pero puede **imprimírsele orientación** mediante la **inicialización** de esa simulación en texto (el prompt) y así materializar diferentes "simulacros", los cuales sí pueden tener objetivos dentro de la simulación.

Problema: datos de entrenamiento limitados y no fidedignos.




El espacio de interpolación


A minimal language model



El espacio de interpolación



Jagged Frontier of AI Capabilities




Inside frontier Expect hard for AI	Outside the frontier Expect easy for AI
Speak any language (mostly) grammatically	Label grammar terms (metalinguage)
Explain $e = mc^2$	Multiply numbers
Write a poem	Reverse a random string of letters
Count people in a story	Count words in a paragraph
Generate complex photos	Place things to the right of other things

One
word
leads

Escalado
masivo del
entrenamiento



Large Language Model (LLM) Sizes Over Time (2018-2026)



Supervisión y alineamiento

Premise

Russian cosmonaut Valery Polyakov set the record for the longest amount of time spent in space.

Hypothesis

Russians hold the record for the longest stay in space.

Target

Entailment
Not entailment



Options:

- yes
- no

Template 1

Russian Cosmonaut Valery Polyakov set the record for the longest amount of time spent in space.

Based on the paragraph above, can we conclude that

Russians hold the record for the longest stay in space?

OPTIONS

- yes
- no

Template 2

Read the following and determine if the hypothesis can be inferred from the premise:

Premise: **<premise>**

Hypothesis: **<hypothesis>**

<options>

Template 3, ...



GPT-3 175B zero shot




GPT-3 175B few-shot



FLAN 137B zero-shot

Performance
on unseen
task types



Supervisión y alineamiento

¡Hola! Muy bien, gracias. ¿Y tú, cómo estás? ¿En qué puedo ayudarte hoy?

📄 ↗️ ↻️ 😊 🗣️ ...

Hola, ¿cómo estás?



¿Qué hora es?

```
[
  {"role": "system", "content": "Eres un asistente útil."},
  {"role": "user", "content": "Hola, ¿cómo estás?"},
  {"role": "assistant", "content": "¡Hola! Muy bien, gracias. ¿Y tú, cómo estás? ¿En qué puedo ayudarte hoy?"},
  {"role": "user", "content": "¿Qué hora es?"}
]
```

```
<|im_start|>system
Eres un asistente útil.<|im_end|>
<|im_start|>user
Hola, ¿cómo estás?<|im_end|>
<|im_start|>assistant
¡Hola! Muy bien, gracias. ¿Y tú, cómo estás? ¿En qué puedo ayudarte hoy?<|im_end|>
<|im_start|>user
¿Qué hora es?<|im_end|>
```

Supervisión y alineamiento


Step 1

Collect demonstration data and train a supervised policy.

A prompt is sampled from our prompt dataset.

A labeler demonstrates the desired output behavior.

This data is used to fine-tune GPT-3.5 with supervised learning.




Step 2

Collect comparison data and train a reward model.

A prompt and several model outputs are sampled.

A labeler ranks the outputs from best to worst.

This data is used to train our reward model.



Step 3

Optimize a policy against the reward model using the PPO reinforcement learning algorithm.


A new prompt is sampled from the dataset.

The PPO model is initialized from the supervised policy.


The policy generates an output.

The reward model calculates a reward for the output.

The reward is used to update the policy using PPO.



Supervisión y
alineamiento



Capacidades de desarrollo de software de un LLM

```
***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
```

```
***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
4     total += n
```

```
1 def sum_list(data):
2     return sum(data)
```

```
***
1 def average(data):
2     if not data:
3         return 0
4     return sum(data) / len(data)
```

```
***
1 def average(data):
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
```

```
***
1 def average(data: list[float]) -> float:
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
```

```
***
1 from typing import Sequence
2
3 def average(data: Sequence[float]) -> float:
4     if not data:
5         raise ValueError(
6             "Data must not be empty"
7         )
8     return sum(data) / len(data)
```



Kata Training

8 kyu Even or Odd ✓

☆ 2696 🏆 493 📈 91% of 34,796 👤 147,433 of 509,399 👤 suuuzi

Instructions Output Past Solutions

Create a function that takes an integer as an argument and returns "Even" for even numbers or "Odd" for odd numbers.

MATHEMATICS FUNDAMENTALS

 **GitLab**
GitLab is the most comprehensive AI-powered

Built on Qualified

Python

3.11

VIM EMACS

Solution

```
1 def even_or_odd(number):
2     pass
```

Sample Tests

```
1 import codewars_test as test
2 from solution import even_or_odd
3
4 @test.describe("Example tests")
5 def example_tests():
6
7     @test.it('even_or_odd(1) should return "Odd"')
8     def test_positive_odd():
9         test.assert_equals(even_or_odd(1), "Odd")
```

SKIP

VIEW SOLUTIONS

DISCUSS (269)

RESET

TEST

ATTEMPT

```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:

```

```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
4     total += n

```

Rendimiento humano en Codewars

```

1 def sum_list(data):
2     return sum(data)
#

```

```

***
1 def average(data):
2     if not data:
3         return 0
4     return sum(data) / len(data)
#

```

```

***
1 def average(data):
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
#

```

```


1 def average(data: list[float]) -> float:
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
#

```

```

***
1 from typing import Sequence
2
3 def average(data: Sequence[float]) -> float:
4     if not data:
5         raise ValueError(
6             "Data must not be empty"
7         )
8     return sum(data) / len(data)
#

```



```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:

```

```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
4     total += n

```

Capacidades de desarrollo de software de un LLM

```

1 def sum_list(data):
2     return sum(data)

```

```

***
1 def average(data):
2     if not data:
3         return 0
4     return sum(data) / len(data)

```

```

***
1 def average(data):
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)

```

```


***
1 def average(data: list[float]) -> float:
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)

```

```

***
1 from typing import Sequence
2
3 def average(data: Sequence[float]) -> float:
4     if not data:
5         raise ValueError(
6             "Data must not be empty"
7         )
8     return sum(data) / len(data)

```



```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:

```

```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
4     total += n

```

Capacidades de desarrollo de software de un LLM

```

1 def sum_list(data):
2     return sum(data)
3

```

```

***
1 def average(data):
2     if not data:
3         return 0
4     return sum(data) / len(data)

```

```

***
1 def average(data):
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
7

```

```


***
1 def average(data: list[float]) -> float:
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
7

```

```

***
1 from typing import Sequence
2
3 def average(data: Sequence[float]) -> float:
4     if not data:
5         raise ValueError(
6             "Data must not be empty"
7         )
8     return sum(data) / len(data)
9

```



```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:

```

```

***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
4     total += n

```

Capacidades de desarrollo de software de un LLM

```

1 def sum_list(data):
2     return sum(data)
3

```

```

***
1 def average(data):
2     if not data:
3         return 0
4     return sum(data) / len(data)

```

```

***
1 def average(data):
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
7

```

```


***
1 def average(data: list[float]) -> float:
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
7

```


```

***
1 from typing import Sequence
2
3 def average(data: Sequence[float]) -> float:
4     if not data:
5         raise ValueError(
6             "Data must not be empty"
7         )
8     return sum(data) / len(data)
9

```



Capacidades de desarrollo de software de un LLM





Cadenas de pensamiento

Standard Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: The answer is 11.

Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The answer is 27. ❌

Chain-of-Thought Prompting

Model Input

Q: Roger has 5 tennis balls. He buys 2 more cans of tennis balls. Each can has 3 tennis balls. How many tennis balls does he have now?

A: Roger started with 5 balls. 2 cans of 3 tennis balls each is 6 tennis balls. $5 + 6 = 11$. The answer is 11.


Q: The cafeteria had 23 apples. If they used 20 to make lunch and bought 6 more, how many apples do they have?

Model Output

A: The cafeteria had 23 apples originally. They used 20 to make lunch. So they had $23 - 20 = 3$. They bought 6 more apples, so they have $3 + 6 = 9$. The answer is 9. ✅

Un LLM entrenado con cadenas de pensamiento puede simular procesos complejos paso a paso, e interpolar y reutilizar estos pasos intermedios.

Cadenas de pensamiento



BaLLMatro



BaLLMatro



A vertical decorative graphic on the left side of the page. It features several playing cards floating and falling. At the top, a 4 of Spades and a 6 of Diamonds are visible. Below them, a 3 of Clubs and a 5 of Hearts are shown. Further down, a 9 of Clubs and a 10 of Hearts are depicted. At the bottom, a large Ace of Hearts is prominent. The cards are surrounded by glowing red and white particle effects, giving a sense of motion.

BaLLMatro





BaLLMatro


5 Chips x 1 Mult	High Card	When no other hand is possible, the one highest card in your hand. Aces are counted high for this hand.	
10 Chips x 2 Mult	Pair	Two cards with a matching rank. Suits may differ.	
20 Chips x 2 Mult	Two Pair	Two cards with a matching rank, and two cards with any other matching rank. Suits may differ.	
30 Chips x 3 Mult	Three of a Kind	Three cards with a matching rank. Suits may differ.	
30 Chips x 4 Mult	Straight	Five cards in consecutive order which are not all from the same suit. Aces can be counted high or low, but not both at once.	
30 Chips x 4 Mult	Straight	Five cards in consecutive order which are not all from the same suit. Aces can be counted high or low, but not both at once.	
35 Chips x 4 Mult	Flush	Five cards of any rank, all from a single suit.	
40 Chips x 4 Mult	Full House	Three cards with a matching rank, and two cards with any other matching rank, with cards from two or more suits.	
50 Chips x 7 Mult	Four of a Kind	Four cards with a matching rank. Suits may differ.	
100 Chips x 8 Mult	Straight Flush	Five cards in consecutive order, all from a single suit.	

Nr	Joker	Effect	Cost	Rarity	Unlock Requirement	Type	Act
1	 Joker	+4 Mult	\$2	Common	Available from start.	+m	Indep.
2	 Greedy Joker	Played cards with Diamond suit give +3 Mult when scored	\$5	Common	Available from start.	+m	On Scored
3	 Lusty Joker	Played cards with Heart suit give +3 Mult when scored	\$5	Common	Available from start.	+m	On Scored
4	 Wrathful Joker	Played cards with Spade suit give +3 Mult when scored	\$5	Common	Available from start.	+m	On Scored
5	 Gluttonous Joker	Played cards with Club suit give +3 Mult when scored	\$5	Common	Available from start.	+m	On Scored
6	 Jolly Joker	+8 Mult if played hand contains a Pair	\$3	Common	Available from start.	+m	Indep.

BaLLMatro



BaLLMatro



En ciertos contextos podemos medir de forma objetiva y automática la calidad de una respuesta de LLM. Esto es posible en tareas como matemáticas, programación, preguntas de opción múltiple...

Refuerzo sobre tareas verificables

```
def incr_list(l: list):
    """Return list with elements incremented by 1.
    >>> incr_list([1, 2, 3])
    [2, 3, 4]
    >>> incr_list([5, 3, 5, 2, 3, 3, 9, 0, 123])
    [6, 4, 6, 3, 4, 4, 10, 1, 124]
    """
    return [i + 1 for i in l]
```

```
def solution(lst):
    """Given a non-empty list of integers, return the sum of all of the odd elements
    that are in even positions.

    Examples
    solution([5, 8, 7, 1]) ==>12
    solution([3, 3, 3, 3, 3]) ==>9
    solution([30, 13, 24, 321]) ==>0
    """
    return sum(lst[i] for i in range(0, len(lst)) if i % 2 == 0 and lst[i] % 2 == 1)
```

Professional Law

As Seller, an encyclopedia salesman, approached the grounds on which Hermit's house was situated, he saw a sign that said, "No salesmen. Trespassers will be prosecuted. Proceed at your own risk."

Although Seller had not been invited to enter, he ignored the sign and drove up the driveway toward the house. As he rounded a curve, a powerful explosive charge buried in the driveway exploded, and Seller was injured. Can Seller recover damages from Hermit for his injuries?

- (A) Yes, unless Hermit, when he planted the charge, intended only to deter, not harm, intruders. ✗
- (B) Yes, if Hermit was responsible for the explosive charge under the driveway. ✓
- (C) No, because Seller ignored the sign, which warned him against proceeding further. ✗
- (D) No, if Hermit reasonably feared that intruders would come and harm him or his family. ✗

Question and choices

Methylcyclopentadiene was allowed to react with methyl isoamyl ketone and a catalytic amount of pyrrolidine. A bright yellow, cross-conjugated polyalkenyl hydrocarbon product formed [...]. How many chemically distinct isomers make up the final product (not counting stereoisomers)?

- (a) 2
- (b) 16
- (c) 8
- (d) 4

Correct answer (b)

Explanation Methylcyclopentadiene exists as an interconverting mixture of 3 isomers [...] if there are 4 dienes, and 4 different directions of approach the dienophile can take to each of them, there are $4 \times 4 = 16$ possible products.

MATH Dataset (Ours)

Problem: Tom has a red marble, a green marble, a blue marble, and three identical yellow marbles. How many different groups of two marbles can Tom choose?

Solution: There are two cases here: either Tom chooses two yellow marbles (1 result), or he chooses two marbles of different colors ($\binom{4}{2} = 6$ results). The total number of distinct pairs of marbles Tom can choose is $1 + 6 = 7$.

Problem: If $\sum_{n=0}^{\infty} \cos^{2n} \theta = 5$, what is $\cos 2\theta$?
Solution: This geometric series is $1 + \cos^2 \theta + \cos^4 \theta + \dots = \frac{1}{1 - \cos^2 \theta} = 5$. Hence,

$$\cos^2 \theta = \frac{4}{5}. \text{ Then } \cos 2\theta = 2 \cos^2 \theta - 1 = \frac{3}{5}.$$

Problem: The equation $x^2 + 2x = i$ has two complex solutions. Determine the product of their real parts.

Solution: Complete the square by adding 1 to each side. Then $(x+1)^2 = 1+i = e^{\frac{\pi}{4}} \sqrt{2}$, so $x+1 = \pm e^{\frac{\pi}{8}} \sqrt[4]{2}$. The desired product is then $(-1 + \cos(\frac{\pi}{8}) \sqrt[4]{2})(-1 - \cos(\frac{\pi}{8}) \sqrt[4]{2}) =$

$$1 - \cos^2(\frac{\pi}{8}) \sqrt{2} = 1 - \frac{(1 + \cos(\frac{\pi}{4}))}{2} \sqrt{2} = \frac{1 - \sqrt{2}}{2}.$$

Hendrycks et al - Measuring Mathematical Problem Solving With the MATH Dataset




Chen et al - Evaluating Large Language Models Trained on Code

Hendrycks et al - Measuring Massive Multitask Language Understanding

Rein et al - GPQA: A Graduate-Level Google-Proof Q&A Benchmark

Aprendizaje por refuerzo: usar los verificadores basados en reglas como función de recompensa y aplicar algoritmos de aprendizaje por refuerzo conocidos, como PPO, para aumentar las probabilidades de que el LLM genere respuestas con puntuaciones altas.

Refuerzo
sobre tareas
verificables



Uso de herramientas y agencia



(1) Hotspot QA

Question: Aside from the Apple Remote, what other device can control the program Apple Remote was originally designed to interact with?

(1a) Standard

Answer: iPod



(1b) CoT (Reason Only)

Thought: Let's think step by step. Apple Remote was originally designed to interact with Apple TV. Apple TV can be controlled by iPhone, iPad, and iPod Touch. So the answer is iPhone, iPad, and iPod Touch.

Answer: iPhone, iPad, iPod Touch



(1c) Act-Only

Act 1: Search[Apple Remote]
Obs 1: The Apple Remote is a remote control ...

Act 2: Search[Front Row]
Obs 2: Could not find [Front Row]. Similar: ...

Act 3: Search[Front Row (software)]

Obs 3: Front Row is a discontinued media center software ...

Act 4: Finish[yes]



(1d) ReAct (Reason + Act)

Thought 1: I need to search Apple Remote and find the program it was originally designed to interact with.

Act 1: Search[Apple Remote]

Obs 1: The Apple Remote is a remote control introduced in October 2005 by Apple ... originally designed to control the Front Row media center program ...

Thought 2: Apple Remote was originally designed to control the Front Row media center program. I need to search Front Row next and find what other device can control it.

Act 2: Search[Front Row]

Obs 2: Could not find [Front Row]. Similar: ['Front Row Seat to Earth', 'Front Row Motorsports', 'Front Row (software)', ...]

Thought 3: Front Row is not found. I need to search Front Row (software) .

Act 3: Search[Front Row (software)]

Obs 3: Front Row is a discontinued media center software ...

Thought 4: Front Row (software) is controlled by an Apple Remote or the keyboard function keys. So the answer is keyboard function keys.

Act 4: Finish[keyboard function keys]




Un LLM entrenado con cadenas de pensamiento y uso de herramientas adquiere capacidades agénticas, ya que puede “razonar” cuándo debe usar una herramienta, la cual es invocada por un arnés externo.

Unido al aprendizaje por refuerzo en tareas verificables, el entrenamiento sesga al LLM a explorar su entorno mediante el uso de herramientas, con el fin de alcanzar el objetivo establecido en sus condiciones iniciales (prompt).

Uso de
herramientas
y agencia



Share of Frontier Lab Compute by Workload: Bandwidth-Bound Work Now Dominates





Capacidades de desarrollo de software de un agente

```
***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
```

```
***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
4     total += n
```

```
1 def sum_list(data):
2     return sum(data)
```

```
***
1 def average(data):
2     if not data:
```



Capacidades de desarrollo de software de un agente

```
***
1 data = [1, 2, 3, 4]
2 total = 0
3 for n in data:
```

```
4     total += n
```

```
1 def sum_list(data):
2     return sum(data)
```

```
***
1 def average(data):
2     if not data:
3         return 0
4     return sum(data) / len(data)
```

```
***
1 def average(data):
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
```


```
***
1 def average(data: list[float]) -> float:
2     if not data:
3         raise ValueError(
4             "Data must not be empty"
5         )
6     return sum(data) / len(data)
```

```
***
1 from typing import Sequence
2
3 def average(data: Sequence[float]) -> float:
4     if not data:
5         raise ValueError(
6             "Data must not be empty"
7         )
8     return sum(data) / len(data)
```


Successful experiments by language and kata difficulty




Codewars score by programming language



Token usage distribution for successful solutions by kata publication year (10th–90th percentile)





ARC-AGI 3

ARC-AGI-3 LEADERBOARD





El espacio de interpolación del lenguaje

¿Cómo de grande es el “espacio del lenguaje”?

- LLM moderno: ~200K tokens de vocabulario, ~1M de tokens de contexto.
- Espacio combinatorial: $200000^{1000000}$ posibles textos.
- Corpus de entrenamiento no-supervisado: **36** trillion tokens ($3.6 \cdot 10^{13}$).
- Estimación de tokens de SFT + refuerzo: ~ **36** trillion tokens ($3.6 \cdot 10^{13}$).
- Fracción del espacio del texto observada $\leq \frac{2 \cdot 3.6 \cdot 10^{13}}{200000^{1000000}} \simeq 7.27 \cdot 10^{-5301017} \simeq 0$

¡Pero el lenguaje tiene una estructura especial!

- **Sintaxis**: no todas las combinaciones de tokens son factibles.
- **No-uniformidad**: la distribución de probabilidad del lenguaje es altamente no uniforme.
- **Composicionalidad**: existen subestructuras que se repiten y se componen.



Esto parece permitir al LLM generar simulacros coherentes con alta probabilidad, e interpolar entre conceptos cercanos en el espacio del lenguaje, **incluso realizando conexiones no conocidas hasta el momento.**


Vacíos en el
espacio
interpolado



Genera una imagen fotorrealista de Ada Lovelace y Alan Turing haciendo una batalla de farmear aura enfrente de la Escuela Politécnica Superior de la Universidad Autónoma de Madrid. Te paso una foto del edificio para que la tomes como referencia. Investiga en internet cómo son las batallas de farmear aura, para que entiendas en qué consisten y puedas representarla fielmente.




Vacíos en el
espacio
interpolado



David Bessis



"The **Overhang** consists of the *unrealized capital gains of past mathematical creativity*, the latent value from connecting the dots in the existing corpus. [...]

Mathematician X states problem A, mathematician Y crafts concept B, then mathematician Z notices that B trivially solves A".




"I expect the Overhang to be absolutely gigantic."

Vacíos en el
espacio
interpolado




AI


Llenado del
espacio
interpolado




Perplexity of generated data points evaluated by
model trained with real wikitext2



Llenado del
espacio
interpolado



Verificación y
canonización



Thank you